

# Training Documentation

## Traffic Sign Classification using ResNet-18 on GTSRB Dataset

---

### Project Overview:

This project involves training a **ResNet-18** model to classify German traffic signs using the **GTSRB (German Traffic Sign Recognition Benchmark)** dataset.

The system processes traffic sign images, learns patterns, and predicts the correct traffic sign category with high accuracy.

---

### Dataset Information:

- **Dataset Name:** GTSRB - German Traffic Sign Recognition Benchmark
  - **Files Used:**
    - Train.csv — Training data with image paths and labels
    - Test.csv — Validation data with image paths and labels
  - **Image Labels:** Traffic sign class IDs ranging from 0 to 42 (43 classes total).
- 

### Project Dependencies:

- **Libraries Used:**
    - torch
    - torchvision
    - PIL (Pillow)
    - pandas
  - **Pretrained Model:**
    - ResNet-18 (from torchvision.models)
- 

### Training Steps Explained:

---

#### 1. Device Configuration

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

- Checks for GPU availability (CUDA).
- Falls back to CPU if a GPU is not found.

---

## 2. Image Transformations

```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std=[0.229, 0.224, 0.225]),
])
```

- **Resize:** All images are resized to 224×224 pixels (ResNet-18 input requirement).
- **Normalization:** Images are normalized with ImageNet dataset means and standard deviations (since the pretrained ResNet-18 was trained on ImageNet).

---

## 3. Custom Dataset Class

```
class TrafficDataset(Dataset):
```

```
...
```

- Custom PyTorch Dataset to load the GTSRB dataset.
- Reads the CSV files (Train.csv, Test.csv), loads corresponding images, and applies transformations.
- Returns a tuple (image\_tensor, label) for each sample.

---

## 4. Data Loading

```
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
```

```
valid_loader = DataLoader(valid_dataset, batch_size=32, shuffle=False)
```

- **Batch size:** 32 images per batch.
- **Shuffle:** Enabled for training data to improve generalization.

---

## 5. Model Preparation

```
model = models.resnet18(pretrained=True)
```

```
num_classes = len(set(train_dataset.data['ClassId']))
```

```
model.fc = nn.Linear(model.fc.in_features, num_classes)
```

```
model = model.to(device)
```

- **Base Model:** Pretrained ResNet-18.

- **Modification:** Final fully connected layer (fc) adjusted for 43 output classes (instead of default 1000).
  - **Transfer Learning:** We leverage pretrained features and fine-tune the model for traffic signs.
- 

## 6. Loss Function & Optimizer

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.Adam(model.parameters(), lr=1e-4)
```

- **Loss Function:** Cross Entropy Loss (standard for multi-class classification).
  - **Optimizer:** Adam Optimizer with a learning rate of 0.0001.
- 

## 7. Training Loop

```
for epoch in range(3):
```

```
...
```

- **Epochs:** 3 (each epoch represents a full pass over the training set).
- **Steps Per Epoch:**
  - Set model to training mode.
  - Forward pass: Images → Model → Predictions.
  - Loss computation.
  - Backward pass: Gradients are computed.
  - Parameters updated via the optimizer.
  - Metrics collected: Running Loss, Accuracy.

### Training Results:

- **Epoch 1:** Loss: 0.1907 | Accuracy: 95.86%
  - **Epoch 2:** Loss: 0.0089 | Accuracy: 99.82%
  - **Epoch 3:** Loss: 0.0077 | Accuracy: 99.82%
- 

## 8. Validation Loop

```
model.eval()
```

```
with torch.no_grad():
```

```
...
```

- Model is switched to **evaluation mode** (no dropout, no batch norm update).

- No gradient calculation (`torch.no_grad()` improves speed and reduces memory usage).
- Predictions are made on the validation dataset.

#### Validation Result:

- **Validation Accuracy:** 99.14%
- 

#### Final Model Performance:

Phase	Accuracy
Training (Epoch 3)	99.82%
Validation	99.14%

---

#### Key Highlights:

- **Pretrained Model:** Leveraging ResNet-18 significantly boosts training speed and accuracy.
  - **Custom Dataset Loader:** Flexibility to work with CSV file structures.
  - **Excellent Accuracy:** >99% accuracy in just 3 epochs due to fine-tuning on powerful ResNet features.
- 

#### Potential Improvements:

- Fine-tune more layers of ResNet-18 (not just fc layer).
  - Use data augmentation (random rotation, horizontal flip, color jitter) for better robustness.
  - Try advanced schedulers like CosineAnnealingLR for dynamic learning rate adjustments.
  - Train for more epochs for even finer learning.
  - Save the best model using `torch.save(model.state_dict(), 'traffic_sign_model.pth')`.
-